

Learning to Reason in 13 Parameters

John X. Morris^{1,2}, Niloofar Mireshghallah^{1,3}, Mark Ibrahim¹, Saeed Mahloujifar¹

¹FAIR at Meta, ²Cornell University, ³Carnegie Mellon University

Recent research has shown that language models can learn to *reason*, often via reinforcement learning. Some work even trains low-rank parameterizations for reasoning, but conventional LoRA cannot scale below the model dimension. We question whether even rank=1 LoRA is necessary for learning to reason and propose TinyLoRA, a method for scaling low-rank adapters to sizes as small as one parameter. Within our new parameterization, we are able to train the 8B parameter size of Qwen2.5 to 91% accuracy on GSM8K with only 13 trained parameters in bf16 (26 total bytes). We find this trend holds in general: we are able to recover 90% of performance improvements while training 1000x fewer parameters across a suite of more difficult learning-to-reason benchmarks such as AIME, AMC, and MATH500. Notably, we are only able to achieve such strong performance with RL: models trained using SFT require 100 – 1000x larger updates to reach the same performance.

Date: February 5, 2026



1 Introduction

Modern language models are post-trained to "reason" in additional thinking tokens through reinforcement learning. Training for reasoning typically takes place through reinforcement learning with verifiable rewards (RLVR) and evaluated in reasoning-intensive domains such as math and coding (OpenAI et al., 2024; Shao et al., 2024; Luo et al., 2025).

Although many practitioners finetune the entire model for reasoning (typically in the billions of parameters) (Guha et al., 2025), a popular class of parameter-efficient methods such as Low-Rank Adaptation (LoRA) (Hu et al., 2021) reduce this requirement from billions to just millions of "adapter" parameters, a reduction of several orders of magnitude. Even so, model updates remain large: for Llama3-8B, running LoRA at its smallest setting (rank 1) finetunes a minimum of 3M parameters (Schulman and Lab, 2025). This parameter count looks surprisingly large compared to past work (Cuccu et al., 2019), which showed certain toy settings such as Atari games can be adequately solved using just six neurons.

There are many advantages to training with fewer parameters. Finetuning fewer parameters reduces per-GPU memory usage as well as the communication cost of distributed training. At inference time, more LoRAs can be stored in memory: a 10x reduction in the size of an adapter allows 10x more LoRAs to be served concurrently (Chen et al., 2023), enabling efficient personalization at scale. Past work (Biderman et al., 2024) also shows that low-rank model updates can mitigate forgetting by retaining important information present in the base model.

Given these benefits, how can we achieve even greater parameter efficiency? Prior work has focused primarily on supervised finetuning (Hu et al., 2021; Li and Liang, 2021; Zaken et al., 2022; Biderman et al., 2024). We observe that good performance in supervised learning requires absorbing many more bits of information into the model than in RL. Excitingly, the converse also holds: when finetuning with RL, we can achieve comparable performance with much smaller updates. This observation leads us to propose an RL-based alternative that enables finetuning with orders of magnitude fewer parameters than standard LoRA-based SFT.

We consider conventional LoRA (Hu et al., 2021) and the smaller LoRA-XS (Hu et al., 2021) which allow us to scale update sizes from 10^6 down to 10^4 parameters. On math tasks, RL is much more effective than SFT in the low-parameter regime (under 1 million parameters trained). On the GSM8K dataset of math word problems, Qwen-2.5-7B-Instruct improves from 76% to 95% while training just 10,000 parameters.




Figure 1 Using Qwen2.5-7B-Instruct as a base model, our TinyLoRA achieves performance within 5% of full finetuning on GSM8K with only 13 parameters. Dashed lines indicate untrained and full-FT baselines.




Figure 2 Using Qwen2.5-7B-Instruct as a base model, SFT works best with larger update sizes of at least 1M parameters.

Since models experience such little performance degradation at 10^4 parameters, we develop a new method (TinyLoRA) to scale LoRA down arbitrarily: our smallest update size is just a single parameter. TinyLoRA is an extra low-rank variant of LoRA that scales smoothly down to just one trained parameter. When training Qwen2.5-7B-Instruct with TinyLoRA and GRPO on GSM8K, **we achieve 91% accuracy while finetuning just 13 total parameters**, an update size of just 26 bytes.

We show that such low-capacity model updates only succeed in the case of RL, and are not information-dense enough to work well for SFT. On GSM8K, models trained with GRPO can obtain 90% accuracy with fewer than 100 parameters, while models of this capacity trained with SFT barely outperform base models on the training set. These findings imply that RL makes fundamentally more information-dense updates than SFT and is more suitable for training in the tiny update size regime.

We also evaluate TinyLoRA on more sophisticated learning-to-reason benchmarks including MATH, AIME, and AMC. We again observe strong performance in the extremely-low parameter regime. As one datapoint, finetuning Qwen2.5-7B-Instruct for just 196 parameters retains 87% of its absolute performance improvement, averaged across six difficult math benchmarks.

2 Related Work

Parameter-efficient finetuning. Our progress scaling LoRA to low parameter counts is part of an ongoing effort to improve low-rank adaptation as evidenced through the slew of proposed related methods such as VeRA (Kopiczko et al., 2023), VB-LoRA (Li et al., 2024b), LoRA-XS (Bałazy et al., 2025), UniLoRA (Li et al., 2025), AdaLoRA (Zhang et al., 2023), LoRA Drop (Zhou et al., 2024), NoRA (Lin et al., 2024), WeightLoRA (Veprikov et al., 2025), and ShareLoRA (Song et al., 2025). We do not advocate for a specific initialization or even parameterization. Rather, we focus on simple methods for scaling to these techniques to low numbers of parameters ($<10K$), where the vast majority of other work on low-rank adaptation is done at the 10K–10M parameter scale. We show that even more efficient parameterizations are possible because we use larger models (billions of parameters) and train with RL as opposed to SFT.

Learning under data and parameter constraints. Another line of work questions how much data is necessary for reasoning, finding that a thousand examples is sufficient for SFT (Ye et al., 2025) and, under certain settings, RL can learn with only one example (Wang et al., 2025). Our analysis of the pareto frontier between update size and performance relates to the general principle that overparameterized models lie on a manifold of a low *intrinsic dimension* (Li et al., 2018; Aghajanyan et al., 2020).




Figure 3 Minimal-sized parameter update to hit threshold of maximum performance vs. backbone model size. Larger models require smaller updates to reach e.g. 95% of peak performance.




Figure 4 Performance ablation using Qwen2.5-3B-Instruct under extremely small update size (<1KB). Surprisingly, storing parameters in fp32 is most performant bit-for-bit.

Comparing SFT and RL. A line of recent work compares supervised finetuning with reasoning tokens (SFT) to reinforcement learning-based methods such as policy gradient. Our analysis of the relationship between data size and LoRA performance relates to finetuning scaling laws (Zhang et al., 2024a) and studies of LoRA learning capacity (Biderman et al., 2024). Some works argue that LoRA makes only superficial changes to models (Shuttleworth et al., 2025). Other work shows that the updates made by RL are small and relatively sparse (Mukherjee et al., 2025). Recently, a number of papers have shown that RL generalizes better out-of-domain than SFT (Chu et al., 2025), forgets less (Shenfeld et al., 2025), and generally amplifies capabilities present in the base model rather than injecting new ones (Rajani et al., 2025). Concurrent work (Schulman and Lab, 2025) demonstrates LoRA can match full finetuning even at $r = 1$ but does not try parameterizations smaller than $3M$.

3 Update Capacity of SFT and RL

In this section we revisit two popular learning algorithms, supervised finetuning and reinforcement learning, and consider the amount of capacity it might take to train a model to good performance using either. Throughout, we consider a model π_θ that maps inputs $x \in \mathcal{X}$ (in this work, word problems) to outputs $y \in \mathcal{Y}$ (solutions).

We study SFT with next-token prediction loss:

$$\mathcal{L}_{\text{SFT}}(\theta) = -\mathbb{E}_{(x,y) \sim (\mathcal{X}, \mathcal{Y})} \left[\sum_{t=1}^{|y|} \log \pi_\theta(y_t | x, y_{<t}) \right]$$

and a class of reinforcement learning algorithms known as policy gradient:

$$\nabla_\theta J(\theta) = \mathbb{E}_{x \sim \mathcal{X}; y \sim \pi_\theta} \left[\sum_{t=1}^{|y|} \nabla_\theta \log \pi_\theta(y_t | x, y_{<t}) \cdot R(y) \right]$$

Information content of training data. Consider training a model θ on a sample (x, y) . The information content of (x, y) for the model can be seen as the minimum description length of (x, y) under the model (Rissanen, 1978). This quantity is independent of how (x, y) is sampled.

For SFT, (x, y) is sampled from a distribution $(\mathcal{X}, \mathcal{Y})$ for which the expected reward is high. For simplicity, assume the reward is binary and the underlying distribution is $(\mathcal{X}, \mathcal{Y})$ where $\mathcal{Y} \equiv \pi_\theta(\mathcal{X}) \mid r(\mathcal{X}, \pi_\theta(\mathcal{X})) = 1$.

For policy gradient methods, the observed data is a batch of continuations sampled from $(\mathcal{X}, (\mathcal{Y}, \mathcal{R})^k)$ where $\mathcal{Y} \equiv \pi_\theta(\mathcal{X})$ and $\mathcal{R} \equiv r(\mathcal{X}, \mathcal{Y})$. These samples are regenerated at each epoch using the updated model.

RL presents more data but less information. In one sense, the information presented during RL appears greater than SFT: at each epoch and for each prompt x , we observe k fresh continuations with new entropy, while SFT uses only one continuation repeated across epochs. However, most of this information is noise. Without the reward annotation, the continuations $\mathcal{Y}$ contain no useful signal for improving task performance. The relevant information lies entirely in the reward, which is only $k \cdot H(\mathcal{R})$ bits per prompt—bounded by k when the reward is binary.

Signal separation. Despite this sparsity of useful information in RL, the signal is cleanly separated from noise. Reward-relevant features correlate with r , reward-irrelevant features do not. Resampling amplifies this separation—the correlated signal accumulates while uncorrelated variation cancels.

In SFT, however, the training signal is a demonstration y with no reward annotation. The model cannot distinguish which features of y are task-relevant. Without a mechanism to separate signal from noise, SFT must treat all tokens as equally informative, storing both the useful structure and the irrelevant details.

Hypothesis. We hypothesize that SFT is less capable in low-parameter regimes because minimizing its objective requires a model to absorb many bits of information, only a fraction of which are relevant to task performance. RL receives a sparser, cleaner signal, allowing it to learn effectively with less capacity.

4 Parameter-Efficient Finetuning with TinyLoRA

The low capacity requirements of reinforcement learning lead us to question just how small an adapter might be while still providing adequate performance. Recent work (Schulman and Lab, 2025; Yin et al., 2025) has observed that LoRA can be effective at update sizes as small as 10^6 parameters. In this section we build upon prior work LoRA and LoRA-XS to propose TinyLoRA, a method for parameter-efficient updates that scale down to a single trained parameter.

4.0.1 LoRA and LoRA-XS

Low-rank adaptation, or LoRA (Hu et al., 2021), adapts a frozen linear layer $W \in \mathbb{R}^{d \times k}$ with a low-rank update:

$$W' = W + AB$$

where $A \in \mathbb{R}^{d \times r}$ and $B \in \mathbb{R}^{r \times k}$ are trainable, and W remains frozen. The number of trainable parameters scales as $\mathcal{O}(dr)$ per module. Applying LoRA to m modules across n layers yields $\mathcal{O}(nm dr)$ total parameters—typically millions for billion-parameter models (Biderman et al., 2024; Schulman and Lab, 2025). LoRA-XS (Bałazy et al., 2025) reduces the per-module parameter count from $\mathcal{O}(dr)$ to $\mathcal{O}(r^2)$:

$$W' = W + U \Sigma R V^\top$$

where $U \in \mathbb{R}^{d \times r}$, $\Sigma \in \mathbb{R}^{r \times r}$, and $V \in \mathbb{R}^{k \times r}$ are from the truncated SVD of W , and only $R \in \mathbb{R}^{r \times r}$ is trainable. This can be viewed as learning to recombine the dominant singular directions of W , and outperforms randomly-initialized LoRA in practice.

4.0.2 Reducing the number of trainable parameters with TinyLoRA

Since the size of a LoRA is inherently tied to the model width, we propose a new technique for scaling down the number of trained parameters. Our technique allows us to train models with as few as 1 changed parameter.

Method	Trainable Parameters	Minimum
FT	$\mathcal{O}(nmd^2)$	nmd^2
LoRA	$\mathcal{O}(nm d r)$	$2nmd$
LoRA-XS	$\mathcal{O}(nm r^2)$	nm
VeRA	$\mathcal{O}(nm(d + r))$	$2nmd$
TinyLoRA	$\mathcal{O}(nm u)$	1

Table 1 Parameter usage comparison per-layer with m adapted modules per layer, model width d , rank r , and TinyLoRA projection dimension u . With weight tying across modules, TinyLoRA can reduce to a single trainable parameter.

Reducing the size of R . Even with $r = 1$, LoRA-XS requires at least one parameter per adapted module. We reduce this further by replacing the $r \times r$ matrix R with a low-dimensional trainable vector $\mathbf{v} \in \mathbb{R}^u$ projected through a fixed random tensor $P \in \mathbb{R}^{u \times r \times r}$. The update rule for TinyLoRA is:

$$W' = W + U \Sigma \left(\sum_{i=1}^u v_i P_i \right) V^\top$$

where U, Σ, V are from the truncated SVD of W , and $P_i \in \mathbb{R}^{r \times r}$ are fixed random matrices. Each module trains only u parameters. With weight tying across m modules in n layers, the total trainable parameters scale as $\mathcal{O}(nm u / n_{\text{tie}})$, reducing to a single parameter when all modules share weights.

Parameter sharing. Prior work (Detrmers et al., 2023; Biderman et al., 2024) showed that LoRA performs best when applied to both MLP and attention modules. In a typical transformer architecture such as LLaMA-3 (AI@Meta, 2024), LoRA is applied seven times per block: to the query, key, value, and output projections in self-attention, and to the up, down, and gate projections in the MLP.

For a model like LLaMA-3 70B with 80 layers, even the minimal case of $u = 1$ (or $r = 1$ for LoRA-XS) requires $80 \times 7 = 560$ trainable parameters. We reduce parameter count further by sharing the trainable vector $\mathbf{v}$ across modules. We define the weight tying factor n_{tie} as the number of modules sharing a single $\mathbf{v}$, yielding $nm u / n_{\text{tie}}$ total trainable parameters for n layers and m modules per layer. With full weight tying ($n_{\text{tie}} = nm$), all modules share a single $\mathbf{v}$, reducing the total to just u parameters—as few as one. Differences between TinyLoRA and related methods are shown in Table 1.

5 Experiments

5.1 Design

The goal of our experiments is to improve math reasoning in language models while minimizing the number of trained parameters. We evaluate on GSM8K (Cobbe et al., 2021), a dataset of 7,500 math word problems, and on the harder MATH training set (Hendrycks et al., 2021).

We consider model training via either *supervised finetuning* (SFT), i.e. next-token prediction, or *reinforcement learning*, in particular Group Relative Policy Optimization (GRPO) (Shao et al., 2024). We conduct experiments on instruction-tuned language models from the Llama-3 (AI@Meta, 2024) and Qwen-2.5 (Qwen et al., 2025) families.

We run four baselines: full finetuning, LoRA, LoRA-XS, and TinyLoRA. We test ranks $\{1, 8, 64, 256\}$ for all LoRA variants; for TinyLoRA, we consider sharing $\{1, 8, 64, 256\}$ layers. All our RL experiments use exact-match reward.

For GSM8K, we do not use any KL penalty, and we train on the full dataset for three epochs with 4 samples per problem and a batch size of 64. We sample with a maximum generation length of 4096.

For MATH training, we follow the settings of SimpleRL (Zeng et al., 2025), which includes a larger training dataset containing subsets such as GSM8K and MATH. The data is partitioned by difficulty level; we only

train on the hardest difficulty, which consists of 8,523 problems. We generate data with a maximum prompt length of 1024 and response length of 3072 tokens, and use the ‘boxed’ chat template. We keep all other hyperparameter settings of SimpleRL: KL coefficient 0.001, temperature 1.0, batch size 256, and 8 generations per response.

Since changes in update size are known to alter effective learning rate (Biderman et al., 2024; Schulman and Lab, 2025), we sweep over learning rates $\{10^{-7}, 5 \times 10^{-7}, 10^{-6}, 5 \times 10^{-6}, 10^{-5}, 10^{-4}, 2 \times 10^{-4}\}$ and take the top-performing learning rate at each update size, averaged over three random seeds.

Evaluation. For GSM8K-only training, we evaluate on the GSM8K validation set. For MATH experiments, we additionally evaluate on the standard datasets proposed by SimpleRL (Zeng et al., 2025): MATH500 (Hendrycks et al., 2021), Minerva (Lewkowycz et al., 2022), GAOKAO (Zhang et al., 2024b), OlympiadBench (He et al., 2024), CollegeMath (Feynman Innovations, 2024), AIME 24 (Li et al., 2024a), and AMC23 (He et al., 2024). All datasets represent math word problems of varying degrees of difficulty and come from the canonical Qwen-Math evaluation (Qwen et al., 2025).

Implementing TinyLoRA in vLLM. We run all RL experiments within the open-source VERL framework (Sheng et al., 2024), using vLLM (Kwon et al., 2023) for inference. However, vLLM requires custom kernels for LoRA and only supports a minimum LoRA rank of 4 (and no LoRA variants). We circumvent this issue by modifying training to use the *merged* model weights for inference, and the true LoRA model only for the final forward pass. This creates a natural numerical mismatch between training and inference, which we mitigate via truncated importance sampling (Ionides, 2008; Yao et al., 2025). This approach allows us to test novel parameter-efficient methods without implementing each one in separate low-level code.

6 Results

6.1 Initial low-parameter results: TinyLoRA, LoRA-XS, and LoRA

We first train and evaluate Qwen2.5-7B-Instruct on GSM8K with GRPO, with results shown in Figure 1. We observe a relatively smooth interpolation from TinyLoRA (1– 10^3 parameters trained) to LoRA-XS (10^3 – 10^6 parameters trained) to LoRA ($> 10^6$ parameters trained.) We are able to recover 95% of the net performance improvement on GSM8K with only 120 parameters trained. **With TinyLoRA, we see a 4% performance increase when training just a single parameter.**

6.2 SFT vs. RL

We display full sweep results for RL (Figure 1) vs. SFT (Figure 2). We find RL is significantly more efficient at low parameter counts: for a baseline of 76%, models trained with RL score 91% (15% absolute improvement) at 13 parameters and 95% at 120 parameters. Models trained with SFT do much worse: 83% at 13 parameters and 84% at 120.

We also observe a less-smooth transition from LoRA-XS at high rank to LoRA at low rank. We note that SFT training directly on answers is a version of *off-policy* learning, and dynamics may change when training directly on model outputs (Mukherjee et al., 2025; Chen et al., 2025).

6.3 Training on MATH

We investigate the performance of low-parameter adaptation methods under the SimpleRL training framework (Zeng et al., 2025). We train models across update sizes for the 7B Qwen model and evaluate on math reasoning tasks; results are shown in Table 2. We consider updates of $\{100, 10,000, 1M\}$ parameters as compared to full finetuning.

We also visualize metrics across training steps (Figure 5). We visualize the mean reward and mean response length. Larger parameter counts tend to achieve higher rewards and generate longer responses, but all parameter counts are getting some rewards – even update sizes as small as 16 parameters. Visualizing the KL divergence between the log probabilities during training and inference shows negligible KL between the train

#	GSM8K	MATH 500	Minerva Math	Olympiad Bench	AIME24	AMC23	Avg.
Qwen2.5-3B-Instruct							
(0)	76.0	55.0	18.5	21.3	2.1	23.4	32.7
16	80.9	64.0	19.9	23.0	3.0	31.5	37.1
63	85.3	64.1	20.1	26.6	7.3	36.0	39.9
252	85.4	66.4	28.3	29.3	13.3	47.5	45.0
504	86.1	66.6	28.7	30.8	16.7	47.5	46.1
8,064	87.2	67.8	28.3	30.7	10.0	47.5	45.2
129,024	86.7	67.8	29.4	32.3	10.0	55.0	46.9
(3,085,846,528)	87.0	69.0	31.7	33.1	15.0	52.2	48.0
Qwen2.5-7B-Instruct							
(0)	88.2	64.6	25.7	30.1	3.3	30.0	40.3
13	91.8	74.6	27.1	36.3	16.0	54.5	50.1
49	91.5	74.2	26.6	37.2	12.6	55.5	49.6
196	92.2	76.6	37.1	38.8	16.7	57.5	53.2
392	92.2	77.0	35.7	40.1	16.7	65.0	54.4
6,272	91.9	78.0	37.5	41.0	16.7	57.5	53.8
100,352	92.8	78.0	37.1	43.3	16.7	60.0	54.6
(7,615,487,488)	91.7	78.2	38.6	40.4	20.0	62.5	55.2
Qwen2.5-7B-Math							
(0)	65.5	63.6	12.5	25.8	13.3	42.5	37.2
196	72.5	62.0	26.5	31.4	20.0	50.0	43.7
392	86.0	74.8	31.6	37.9	26.7	60.0	52.8
6,272	87.0	77.4	28.7	40.0	16.7	67.5	52.9
100,352	87.0	78.6	32.7	39.0	30.0	62.5	55.0
(7,615,487,488)	90.2	80.2	37.5	39.0	40.0	70.0	59.5

Table 2 Performance on math reasoning using Qwen2.5 models. (0) indicates the base model without training; (*) indicates full finetuning (no LoRA). All experiments use the GRPO algorithm and train for three epochs. Full and baseline results are reported from (Zeng et al., 2025).

and inference models, showing our technique of merging the LoRA weights at each training step is working properly.

6.4 Scaling across model backbones

We then scale TinyLoRA from the XS $r = 1$ case (np parameters) down to a single trained parameter, for Qwen3-8B on GSM8K. Results are shown in Figure 1. For the task of GSM8K, Qwen is more parameter-efficient than LLAMA at every scale. In an absolute sense, Qwen is surprisingly efficient, achieving 94.7% accuracy with only 13 parameters trained. Qwen even achieves decent accuracy at 1 parameter, scoring near 82%. (around a 5% improvement over the baseline). This finding may corroborate recent observations that Qwen in particular may have exposure to similar examples during its pretraining (Wu et al., 2025).

Even in the case of LoRA-XS, which enables adapting layers with a single parameter, we are upper-bounded by the total number of linear layers in the model. Large models can contain hundreds of linear layers, which would allow for hundreds of bytes of storage at $r = 1$. Empirically, we observe that Qwen can still learn nearly as well with only a few hundred parameters.

Although LLAMA has a similar upper-bound (from full finetuning), in this low-parameter setting LLAMA is less responsive than Qwen and reaches 85% with an update size of 1KB (500 parameters trained in bf16).




Figure 5 TinyLoRA performance during training on MATH.




Figure 6 TinyLoRA performance (on GSM8K GRPO) across backbone size, with no-training baselines as dashed lines. Small updates only improve larger models.

Unlike Qwen, when we train fewer than five parameters, LLAMA barely improves performance above baseline.

Ranks 1 to 128. We visualize performance from $r = 1$ to $r = 128$ for Qwen and LLAMA in Figure 1. This covers update sizes from 1KB ($r=1$) to 8MB ($r=128$). We see consistent monotonic increase in performance, suggesting that although the performance boost is decreasing per-parameter, finetuning additional parameters does help up to an update size of around 2MB.

From LoRA-XS to LoRA, across model backbones. We display results with LoRA and with LoRA-XS in Figure 6 across backbone sizes. Our results are mixed: LoRA-XS clearly outperforms LoRA on the smallest model, corroborating findings of (Bałazy et al., 2025). But as model size grows, this trend becomes less pronounced, and degrades to performance scaling with the number of trained parameters. (We note that the number of LoRA modules increases significantly in larger models, and this may make all the difference here.)

In general, low-parameter adaptation works better with larger models – in such cases, smaller adapters appear closer to the performance ceiling. We plot the update size vs. relative performance increase across models from the Qwen2.5-Instruct family in Figure 3. The trend is extremely clear: as model size grows, the model is programmable to within 95% of its full finetuning performance with fewer *absolute* parameters. These trends indicate that extremely large (trillion-scale) models may be easily trainable for many tasks with just a handful of parameters.

6.5 Scaling in the bit-constrained regime

We question how to scale TinyLoRA when the parameter count is small (tens or hundreds of parameters) and the constraint is on the total update size in *bytes*. This scenario is common in distributed training, where communicating model parameter updates can be a major bottleneck. We consider two parameter-sharing strategies: *structured*, where nearby modules of the same type share parameters, and *tiled*, where nearby modules of similar depth share parameters, agnostic of type.

Results across sharing strategies and precisions are shown in Figure 4. Surprisingly, tiling outperforms structured sharing: we see no benefit from sharing parameters between e.g. query projection modules. With all-layer sharing and float16 precision, Qwen still trains to a level of 70% on GSM8K, an absolute improvement of over 10%. We also notice that fp32 outperforms bf16 and float16 even when accounting for its twice-as-large update size in bytes.

7 Ablations

We train a number of smaller models to analyze the effect of TinyLoRA hyperparameters. Using GRPO, we finetune Qwen-2.5-3B-Instruct and LLaMA-3.2-3B-Instruct, sweeping over the frozen SVD rank r , the trainable projection dimension u , and the weight tying factor n_{tie} (the number of modules sharing a single $\mathbf{v}$).

Selection of frozen rank r . We first analyze performance as a function of the frozen rank r (Figure 7). While one might expect that preserving more singular directions would improve expressivity, we find diminishing returns: increasing r from 1 to 2 yields modest gains, but larger values degrade performance. We hypothesize that higher ranks introduce more degrees of freedom in the frozen U , Σ , V components, making optimization of the small trainable vector $\mathbf{v}$ more difficult. Based on these results, we use $r = 2$ for all main experiments.

Trading off u and n_{tie} . For a fixed parameter budget, TinyLoRA requires balancing the per-module projection dimension u against the weight tying factor n_{tie} . Increasing u adds expressivity per module but increases parameter count; increasing n_{tie} reduces parameters but forces modules to share the same $\mathbf{v}$. We plot this tradeoff in Figure 8. Performance generally improves with larger u and smaller n_{tie} —that is, more expressive per-module updates with less sharing. This suggests a clear guideline for practitioners: exhaust the u budget (down to $u = 1$) before increasing n_{tie} .




Figure 7 TinyLoRA performance across frozen rank r and trainable rank $r_{\text{trainable}}$.

8 Discussion

Our work raises interesting questions about how model update size during finetuning scales with model size and by learning objective.




Figure 8 Performance ablation with Qwen2.5-3B-Instruct across numbers of tied layers and trainable ranks.

Scaling trends. Our findings indicate that LoRA becomes *more effective* at smaller parameter counts as model size scales. Given a fixed dataset, larger models can be controlled with fewer parameters. This trend indicates that extremely parameter-efficient methods such as TinyLoRA will only grow in popularity as models continue to scale.

How is this possible? One might ask how it’s possible to learn to solve a difficult task such as GSM8K in as few as 13 parameters. One theory is that the *knowledge* required to solve the task is already stored in the parameters of the model, and only the *style* has to change for task success. In particular, learning to generate longer outputs may be possible in few parameters, and has been shown to greatly improve performance on math and reasoning tasks (Shao et al., 2024).

Peculiarities of Qwen and LLaMa. We also observe that the Qwen-2.5 models are generally much more performant at small update sizes than LLaMA-3. In general, Qwen models of the same family require around 10x fewer parameters updated to reach equivalent performance to LLaMa. This may relate to differences in details in the architecture or pretraining processes of the models.

Limitations of math datasets. We note that our findings are limited to math datasets. Many results seem to indicate that, using reinforcement learning, very small updates (on the order of hundreds of bytes) are enough to facilitate learning to reason with reasonable performance. But these findings are limited to the scope of math-style reasoning, and may or may not generalize to other fields such as science or creative writing. We focus on math as it is the most popular application of reasoning models as of 2025, but look forward to seeing our findings extended and replicated in other domains.

9 Conclusion

We propose TinyLoRA and show that effective models can be tuned with far fewer parameters than previously thought to be necessary. Updates from TinyLoRA often come close to the performance of full finetuning can be expressed in very small file sizes – often less than 1KB in total. Our results show that some models can learn to do certain tasks while absorbing very small numbers of bytes, which calls into question what is actually learned during RL with verification. Finally, we show that this is only possible in the case of RL on large models, and does not work as well when not doing RL or not using large models.

10 Acknowledgements

Thanks to Jessy Lin, Dan Biderman, Sasha Rush, John Schulman, and Kamalika Chaudhuri for helpful feedback on this work.

References

- Armen Aghajanyan, Luke Zettlemoyer, and Sonal Gupta. Intrinsic dimensionality explains the effectiveness of language model fine-tuning, 2020. URL <https://arxiv.org/abs/2012.13255>.
- AI@Meta. Llama 3 model card. 2024. URL https://github.com/meta-llama/llama3/blob/main/MODEL_CARD.md.
- Klaudia Bałazy, Mohammadreza Banaei, Karl Aberer, and Jacek Tabor. Lora-xs: Low-rank adaptation with extremely small number of parameters, 2025. URL <https://arxiv.org/abs/2405.17604>.
- Dan Biderman, Jacob Portes, Jose Javier Gonzalez Ortiz, Mansheej Paul, Philip Greengard, Connor Jennings, Daniel King, Sam Havens, Vitaliy Chiley, Jonathan Frankle, Cody Blakeney, and John P. Cunningham. Lora learns less and forgets less, 2024. URL <https://arxiv.org/abs/2405.09673>.
- Howard Chen, Noam Razin, Karthik Narasimhan, and Danqi Chen. Retaining by doing: The role of on-policy data in mitigating forgetting, 2025. URL <https://arxiv.org/abs/2510.18874>.
- Lequn Chen, Zihao Ye, Yongji Wu, Danyang Zhuo, Luis Ceze, and Arvind Krishnamurthy. Punica: Multi-tenant lora serving, 2023. URL <https://arxiv.org/abs/2310.18547>.
- Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang Tong, Saining Xie, Dale Schuurmans, Quoc V. Le, Sergey Levine, and Yi Ma. Sft memorizes, rl generalizes: A comparative study of foundation model post-training, 2025. URL <https://arxiv.org/abs/2501.17161>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021. URL <https://arxiv.org/abs/2110.14168>.
- Giuseppe Cuccu, Julian Togelius, and Philippe Cudre-Mauroux. Playing atari with six neurons, 2019. URL <https://arxiv.org/abs/1806.01363>.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms, 2023. URL <https://arxiv.org/abs/2305.14314>.
- Feynman Innovations. Maths-college. <https://huggingface.co/datasets/ajibawa-2023/Maths-College>, 2024. revision 9eaa32a.
- Etash Guha, Ryan Marten, Sedrick Keh, Negin Raoof, Georgios Smyrnis, Hritik Bansal, Marianna Nezhurina, Jean Mercat, Trung Vu, Zayne Sprague, Ashima Suvarna, Benjamin Feuer, Liangyu Chen, Zaid Khan, Eric Frankel, Sachin Grover, Caroline Choi, Niklas Muennighoff, Shiye Su, Wanbia Zhao, John Yang, Shreyas Pimpalgaonkar, Kartik Sharma, Charlie Cheng-Jie Ji, Yichuan Deng, Sarah Pratt, Vivek Ramanujan, Jon Saad-Falcon, Jeffrey Li, Achal Dave, Alon Albalak, Kushal Arora, Blake Wulfe, Chinmay Hegde, Greg Durrett, Sewoong Oh, Mohit Bansal, Saadia Gabriel, Aditya Grover, Kai-Wei Chang, Vaishaal Shankar, Aaron Gokaslan, Mike A. Merrill, Tatsunori Hashimoto, Yejin Choi, Jenia Jitsev, Reinhard Heckel, Maheswaran Sathiamoorthy, Alexandros G. Dimakis, and Ludwig Schmidt. Openthoughts: Data recipes for reasoning models, 2025. URL <https://arxiv.org/abs/2506.04178>.
- Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, Jie Liu, Lei Qi, Zhiyuan Liu, and Maosong Sun. Olympiadbench: A challenging benchmark for promoting agi with olympiad-level bilingual multimodal scientific problems, 2024. URL <https://arxiv.org/abs/2402.14008>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset, 2021. URL <https://arxiv.org/abs/2103.03874>.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021. URL <https://arxiv.org/abs/2106.09685>.
- Edward L. Ionides. Truncated importance sampling. *Journal of Computational and Graphical Statistics*, 17(2):295–311, 2008. doi: 10.1198/106186008X320456. URL <https://doi.org/10.1198/106186008X320456>.
- Dawid J. Kopiczko, Tijmen Blankevoort, and Yuki M. Asano. Vera: Vector-based random matrix adaptation. *arXiv preprint arXiv:2310.11454*, 2023. URL <https://arxiv.org/abs/2310.11454>. Revised version v2, 16 Jan 2024.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention, 2023. URL <https://arxiv.org/abs/2309.06180>.

- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving quantitative reasoning problems with language models, 2022. URL <https://arxiv.org/abs/2206.14858>.
- Chunyuan Li, Heerad Farkhor, Rosanne Liu, and Jason Yosinski. Measuring the intrinsic dimension of objective landscapes, 2018. URL <https://arxiv.org/abs/1804.08838>.
- Jia Li, Edward Beeching, Lewis Tunstall, Ben Lipkin, Roman Soletskyi, Shengyi Costa Huang, Kashif Rasul, Longhui Yu, Albert Jiang, Ziju Shen, Zihan Qin, Bin Dong, Li Zhou, Yann Fleureau, Guillaume Lample, and Stanislas Polu. Numinamath, 2024a. URL https://github.com/project-numina/aimo-progress-prize/blob/main/report/numina_dataset.pdf.
- Kaiyang Li, Shaobo Han, Qing Su, Wei Li, Zhipeng Cai, and Shihao Ji. Uni-lora: One vector is all you need. *arXiv preprint arXiv:2506.00799*, 2025. URL <https://arxiv.org/abs/2506.00799>.
- Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation, 2021. URL <https://arxiv.org/abs/2101.00190>.
- Yang Li, Shaobo Han, and Shihao Ji. Vb-lora: Extreme parameter efficient fine-tuning with vector banks, 2024b. URL <https://arxiv.org/abs/2405.15179>.
- Cheng Lin, Lujun Li, Dezhi Li, Jie Zou, Wei Xue, and Yike Guo. Nora: Nested low-rank adaptation for efficient fine-tuning large models, 2024. URL <https://arxiv.org/abs/2408.10280>.
- Michael Luo, Sijun Tan, Justin Wong, Xiaoxiang Shi, William Y. Tang, Manan Roongta, Colin Cai, Jeffrey Luo, Li Erran Li, Raluca Ada Popa, and Ion Stoica. Deepscaler: Surpassing o1-preview with a 1.5b model by scaling rl. <https://pretty-radio-b75.notion.site/DeepScaleR-Surpassing-O1-Preview-with-a-1-5B-Model-by-Scaling-RL-19681902c1468005bed8ca303013a4e2>, 2025. Notion Blog.
- Sagnik Mukherjee, Lifan Yuan, Dilek Hakkani-Tur, and Hao Peng. Reinforcement learning finetunes small subnetworks in large language models, 2025. URL <https://arxiv.org/abs/2505.11711>.
- OpenAI, :, Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry ..., Yuchen He, Yuchen Zhang, Yunyun Wang, Zheng Shao, and Zhuohan Li. Openai o1 system card, 2024. URL <https://arxiv.org/abs/2412.16720>.
- Qwen, :, An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.
- Neel Rajani, Aryo Pradipta Gema, Seraphina Goldfarb-Tarrant, and Ivan Titov. Scalpel vs. hammer: Grpo amplifies existing capabilities, sft replaces them, 2025. URL <https://arxiv.org/abs/2507.10616>.
- Jorma Rissanen. Modeling by shortest data description. *Automatica*, 14(5):465–471, 1978.
- John Schulman and Thinking Machines Lab. Lora without regret. *Thinking Machines Lab: Connectionism*, 2025. doi: 10.64434/tml.20250929. <https://thinkingmachines.ai/blog/lora/>.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024. URL <https://arxiv.org/abs/2402.03300>.
- Idan Shenfeld, Jyothish Pari, and Pulkit Agrawal. RL’s razor: Why online reinforcement learning forgets less, 2025. URL <https://arxiv.org/abs/2509.04259>.
- Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. *arXiv preprint arXiv: 2409.19256*, 2024.
- Reece Shuttlesworth, Jacob Andreas, Antonio Torralba, and Pratyusha Sharma. Lora vs full fine-tuning: An illusion of equivalence, 2025. URL <https://arxiv.org/abs/2410.21228>.
- Yurun Song, Junchen Zhao, Ian G. Harris, and Sangeetha Abdu Jyothi. Sharelora: Parameter efficient and robust large language model fine-tuning via shared low-rank adaptation, 2025. URL <https://arxiv.org/abs/2406.10785>.

- Andrey Veprikov, Vladimir Solodkin, Alexander Zyl, Andrey Savchenko, and Aleksandr Beznosikov. Weightlora: Keep only necessary adapters, 2025. URL <https://arxiv.org/abs/2506.02724>.
- Yiping Wang, Qing Yang, Zhiyuan Zeng, Liliang Ren, Liyuan Liu, Baolin Peng, Hao Cheng, Xuehai He, Kuan Wang, Jianfeng Gao, Weizhu Chen, Shuohang Wang, Simon Shaolei Du, and Yelong Shen. Reinforcement learning for reasoning in large language models with one training example, 2025. URL <https://arxiv.org/abs/2504.20571>.
- Mingqi Wu, Zhihao Zhang, Qiaole Dong, Zhiheng Xi, Jun Zhao, Senjie Jin, Xiaoran Fan, Yuhao Zhou, Huijie Lv, Ming Zhang, Yanwei Fu, Qin Liu, Songyang Zhang, and Qi Zhang. Reasoning or memorization? unreliable results of reinforcement learning due to data contamination, 2025. URL <https://arxiv.org/abs/2507.10532>.
- Feng Yao, Liyuan Liu, Dinghuai Zhang, Chengyu Dong, Jingbo Shang, and Jianfeng Gao. Your efficient rl framework secretly brings you off-policy rl training, August 2025. URL <https://fengyao.notion.site/off-policy-rl>.
- Yixin Ye, Zhen Huang, Yang Xiao, Ethan Chern, Shijie Xia, and Pengfei Liu. Limo: Less is more for reasoning, 2025. URL <https://arxiv.org/abs/2502.03387>.
- Qingyu Yin, Yulun Wu, Zhennan Shen, Sunbowen Li, Zhilin Wang, Yanshu Li, Chak Tou Leong, Jiale Kang, and Jinjin Gu. Evaluating parameter efficient methods for rlvr, 2025. URL <https://arxiv.org/abs/2512.23165>.
- Elad Ben Zaken, Shauli Ravfogel, and Yoav Goldberg. Bitfit: Simple parameter-efficient fine-tuning for transformer-based masked language-models, 2022. URL <https://arxiv.org/abs/2106.10199>.
- Weihaio Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. Simplerl-zoo: Investigating and taming zero reinforcement learning for open base models in the wild, 2025. URL <https://arxiv.org/abs/2503.18892>.
- Biao Zhang, Zhongtao Liu, Colin Cherry, and Orhan Firat. When scaling meets llm finetuning: The effect of data, model and finetuning method, 2024a. URL <https://arxiv.org/abs/2402.17193>.
- Qingru Zhang, Minshuo Chen, Alexander Bukharin, Nikos Karampatziakis, Pengcheng He, Yu Cheng, Weizhu Chen, and Tuo Zhao. Adalora: Adaptive budget allocation for parameter-efficient fine-tuning, 2023. URL <https://arxiv.org/abs/2303.10512>.
- Xiaotian Zhang, Chunyang Li, Yi Zong, Zhengyu Ying, Liang He, and Xipeng Qiu. Evaluating the performance of large language models on gaokao benchmark, 2024b. URL <https://arxiv.org/abs/2305.12474>.
- Hongyun Zhou, Xiangyu Lu, Wang Xu, Conghui Zhu, Tiejun Zhao, and Muyun Yang. Lora-drop: Efficient lora parameter pruning based on output evaluation, 2024. URL <https://arxiv.org/abs/2402.07721>.

A Appendix

A.1 Additional ablations




Figure 9 Performance ablation for Qwen2.5-3B-Instruct with different numbers of tied layers and ranks r .

We ran a number of ablations to isolate the right way to ‘spend’ trainable parameters: is it better to have more, shared parameters or fewer-but-unique parameters? We sweep across values of $r_{\text{trainable}}$ for 3B versions

of Qwen-2.5 and LLAMA-3.2 on GSM8K. Surprisingly, we find higher performance at matched parameter count for a *lower* value of $r_{\text{trainable}}$, indicating parameters are better spent on fewer, unique parameters. We also observe again that Qwen far outperforms Llama at a constant update size: when training just 100 parameters, Llama achieves 60% accuracy, while Qwen averages 74% across 3 random seeds.